

```
from IPython.display import Image  
Image(filename="M:\OLA CASE STUDY\ola-electric-scooter-  
1536x858.jpg",width=500)
```



Important Libraries For Analysis

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
from scipy.stats import zscore  
import seaborn as sns  
import warnings  
warnings.filterwarnings("ignore")  
%matplotlib inline
```

Problem Statement:

Recruiting and retaining drivers is seen by industry watchers as a tough battle for Ola. Churn among drivers is high and it's very easy for drivers to stop working for the service on the fly or jump to Uber depending on the rates. As the companies get bigger, the high churn could become a bigger problem. To find new drivers, Ola is casting a wide net, including people who don't have cars for jobs. But this acquisition is really costly. Losing drivers frequently impacts the morale of the organization and acquiring new drivers is more expensive than retaining existing ones. You are working as a data scientist with the Analytics Department of Ola, focused on driver team attrition. You are provided with the monthly information for a segment of drivers for 2019 and 2020 and tasked to predict whether a driver will be leaving the company or not based on their attributes like:

```
df=pd.read_csv("M:\OLA CASE STUDY\ola_driver.csv")
```

Data Exploration and Imputation

df.head()

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level
0	0	01/01/19	1	28.0	0.0	C23	2
1	1	02/01/19	1	28.0	0.0	C23	2
2	2	03/01/19	1	28.0	0.0	C23	2
3	3	11/01/20	2	31.0	0.0	C7	2
4	4	12/01/20	2	31.0	0.0	C7	2

	Income	Dateofjoining	LastWorkingDate	Joining	Designation	Grade
0	57387	24/12/18	NaN		1	1
1	57387	24/12/18	NaN		1	1
2	57387	24/12/18	03/11/19		1	1
3	67016	11/06/20	NaN		2	2
4	67016	11/06/20	NaN		2	2

	Total Business Value	Quarterly Rating
0	2381060	2
1	-665480	2
2	0	2
3	0	1
4	0	1

df.tail()

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level
19099	19099	08/01/20	2788	30.0	0.0	C27	2
19100	19100	09/01/20	2788	30.0	0.0	C27	2
19101	19101	10/01/20	2788	30.0	0.0	C27	2
19102	19102	11/01/20	2788	30.0	0.0	C27	2
19103	19103	12/01/20	2788	30.0	0.0	C27	2

	Income	Dateofjoining	LastWorkingDate	Joining	Designation	Grade
19099	70254	06/08/20	NaN			2
19100	70254	06/08/20	NaN			2

19101	70254	06/08/20	NaN	2
2				
19102	70254	06/08/20	NaN	2
2				
19103	70254	06/08/20	NaN	2
2				

	Total Business Value	Quarterly Rating
19099	740280	3
19100	448370	3
19101	0	2
19102	200420	2
19103	411480	2

df.shape

(19104, 14)

df.dtypes

```

Unnamed: 0      int64
MMM-YY         object
Driver_ID      int64
Age            float64
Gender         float64
City           object
Education_Level int64
Income         int64
Dateofjoining  object
LastWorkingDate object
Joining Designation int64
Grade          int64
Total Business Value int64
Quarterly Rating int64
dtype: object

```

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 19104 entries, 0 to 19103

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	19104 non-null	int64
1	MMM-YY	19104 non-null	object
2	Driver_ID	19104 non-null	int64
3	Age	19043 non-null	float64
4	Gender	19052 non-null	float64
5	City	19104 non-null	object
6	Education_Level	19104 non-null	int64
7	Income	19104 non-null	int64

```

8   Dateofjoining      19104 non-null object
9   LastWorkingDate    1616 non-null object
10  Joining Designation 19104 non-null int64
11  Grade              19104 non-null int64
12  Total Business Value 19104 non-null int64
13  Quarterly Rating    19104 non-null int64
dtypes: float64(2), int64(8), object(4)
memory usage: 2.0+ MB

```

```
df.describe()
```

	Unnamed: 0	Driver_ID	Age	Gender \
count	19104.000000	19104.000000	19043.000000	19052.000000
mean	9551.500000	1415.591133	34.668435	0.418749
std	5514.994107	810.705321	6.257912	0.493367
min	0.000000	1.000000	21.000000	0.000000
25%	4775.750000	710.000000	30.000000	0.000000
50%	9551.500000	1417.000000	34.000000	0.000000
75%	14327.250000	2137.000000	39.000000	1.000000
max	19103.000000	2788.000000	58.000000	1.000000

	Education_Level	Income	Joining Designation
Grade \			
count	19104.000000	19104.000000	19104.000000
mean	1.021671	65652.025126	1.690536
std	0.800167	30914.515344	0.836984
min	0.000000	10747.000000	1.000000
25%	0.000000	42383.000000	1.000000
50%	1.000000	60087.000000	1.000000
75%	2.000000	83969.000000	2.000000
max	2.000000	188418.000000	5.000000

	Total Business Value	Quarterly Rating
count	1.910400e+04	19104.000000
mean	5.716621e+05	2.008899
std	1.128312e+06	1.009832
min	-6.000000e+06	1.000000
25%	0.000000e+00	1.000000
50%	2.500000e+05	2.000000
75%	6.997000e+05	3.000000
max	3.374772e+07	4.000000

```
df.describe(include='object')
```


	MMM-YY	City	Dateofjoining	LastWorkingDate
count	19104	19104	19104	1616
unique	24	29	869	493
top	01/01/19	C20	23/07/15	29/07/20
freq	1022	1008	192	70

```
df.isnull().sum()
```

Unnamed: 0	0
MMM-YY	0
Driver_ID	0
Age	61
Gender	52
City	0
Education_Level	0
Income	0
Dateofjoining	0
LastWorkingDate	17488
Joining Designation	0
Grade	0
Total Business Value	0
Quarterly Rating	0
dtype:	int64

```
sns.boxplot(x='Age',data=df)
plt.xlabel("Age")
plt.ylabel("Frequency")
plt.title("Age Distribution in boxplot")
```

```
Text(0.5, 1.0, 'Age Distribution in boxplot')
```



```
df['Age']=df['Age'].fillna(df['Age'].mean())
```

```
df['Age'].isnull().sum()
```

```
0
```

```
df.isnull().sum()
```

```

Unnamed: 0      0
MMM-YY          0
Driver_ID       0
Age             0
Gender          52
City            0
Education_Level 0
Income          0
Dateofjoining   0
LastWorkingDate 17488
Joining Designation 0
Grade           0
Total Business Value 0
Quarterly Rating 0
dtype: int64

```

```

non_null_genders=df.Gender.dropna().values
non_null_genders

array([0., 0., 0., ..., 0., 0., 0.])

df.Gender.value_counts()

Gender
0.0    11074
1.0     7978
Name: count, dtype: int64

df['Gender']=df['Gender'].apply(lambda
x:np.random.choice(non_null_genders) if pd.isnull(x) else x)

df['Gender'].isnull().sum()

0

df.Gender.value_counts()

Gender
0.0    11103
1.0     8001
Name: count, dtype: int64

```

Feature Engineering

```

df['Hasleftcompany']=df['LastWorkingDate'].apply(lambda x:0 if
pd.isnull(x) else 1)

df['Hasleftcompany'].value_counts()

Hasleftcompany
0    17488
1     1616
Name: count, dtype: int64

df['Dateofjoining']=pd.to_datetime(df['Dateofjoining'],errors='coerce'
)
current_year=pd.Timestamp.today().year
df['Birthyear']=current_year-df['Age']
df['AgeAtJoining']=df['Dateofjoining'].dt.year-df['Birthyear']

df['Birthyear']=df['Birthyear'].astype('int64')

df['AgeAtJoining']=df['AgeAtJoining'].astype('int64')

df.head()

```

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level
0	0	01/01/19	1	28.0	0.0	C23	2

1	1	02/01/19	1	28.0	0.0	C23	2
2	2	03/01/19	1	28.0	0.0	C23	2
3	3	11/01/20	2	31.0	0.0	C7	2
4	4	12/01/20	2	31.0	0.0	C7	2

	Income	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
0	57387	2018-12-24	NaN		1	1	
1	57387	2018-12-24	NaN		1	1	
2	57387	2018-12-24	03/11/19		1	1	
3	67016	2020-11-06	NaN		2	2	
4	67016	2020-11-06	NaN		2	2	

	Total Business Value	Quarterly Rating	Hasleftcompany	Birthyear	\
0	2381060	2	0	1997	
1	-665480	2	0	1997	
2	0	2	1	1997	
3	0	1	0	1994	
4	0	1	0	1994	

	AgeAtJoining
0	21
1	21
2	21
3	26
4	26

```

df['Dateofjoining']=pd.to_datetime(df['Dateofjoining'],errors='coerce')
df['LastWorkingDate']=pd.to_datetime(df['LastWorkingDate'],errors='coerce')
df['Tenure']=df.apply(lambda row:(pd.Timestamp.today()-
row['Dateofjoining']).days
                        if pd.isnull(row['LastWorkingDate']) else
(row['LastWorkingDate']-row['Dateofjoining']).days, axis=1)

max_income=df['Income'].max()
print(f"Maximum income of any Employee is :{max_income}")
min_income=df['Income'].min()
print(f"Minimum income of any Employee is :{min_income}")

```


Maximum income of any Employee is :188418

Minimum income of any Employee is :10747

```
df['Salary_Range']=pd.cut(df['Income'],bins=[10000,50000,100000,150000,200000],labels=['Low Salary','Medium Salary','High Salary','Upper High Salary'])
```

```
df['Salary_Range']
```

```
0      Medium Salary
1      Medium Salary
2      Medium Salary
3      Medium Salary
4      Medium Salary
```

```
...
19099   Medium Salary
19100   Medium Salary
19101   Medium Salary
19102   Medium Salary
19103   Medium Salary
```

Name: Salary_Range, Length: 19104, dtype: category

Categories (4, object): ['Low Salary' < 'Medium Salary' < 'High Salary' < 'Upper High Salary']

```
df['Salary_Range'].value_counts()
```

```
Salary_Range
Medium Salary      9696
Low Salary         6604
High Salary        2664
Upper High Salary   140
Name: count, dtype: int64
```

```
df.head(1)
```

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level
0	0	01/01/19	1	28.0	0.0	C23	2

	Income	Dateofjoining	LastWorkingDate	Joining	Designation	Grade
0	57387	2018-12-24	NaT			1

	Total Business Value	Quarterly Rating	Hasleftcompany
0	2381060	2	0

	AgeAtJoining	Tenure	Salary_Range
0	21	2226	Medium Salary

```
df['Quarterly Rating'].value_counts()
```

```

Quarterly Rating
1      7679
2      5553
3      3895
4      1977
Name: count, dtype: int64

df['Quarterly_Range']=pd.cut(df['Quarterly
Rating'],bins=[1,2,3,4],labels=['Low','Medium','High'])

df['Quarterly_Range'].value_counts()

Quarterly_Range
Low      5553
Medium   3895
High     1977
Name: count, dtype: int64

max_age=df['Age'].max()
min_age=df['Age'].min()
print(f"Maximum Age of any person in the DataSet is : {max_age}")
print(f"Minimum Age of any person in the DataSet is : {min_age}")

Maximum Age of any person in the DataSet is : 58.0
Minimum Age of any person in the DataSet is : 21.0

df['Riders_Age_Category']=pd.cut(df['Age'],bins=[21,25,30,58],labels=[
'Young Riders','Medium Riders','Old Riders'])

df['Riders_Age_Category'].value_counts()

Riders_Age_Category
Old Riders      13820
Medium Riders   4241
Young Riders    1008
Name: count, dtype: int64

df['MMM-YY']=pd.to_datetime(df['MMM-YY'],format="%d/%m/%y")
df=df.sort_values(by=['Driver_ID','MMM-YY'])
def Rating_increasing(rating):
    return (rating>rating.shift(1)).astype(int)

df['Rating_Increased']=df.groupby('Driver_ID')['Quarterly
Rating'].apply(Rating_increasing).reset_index(level = 0, drop = True)

df[['MMM-YY','Driver_ID','Quarterly Rating','Rating_Increased']]

```

	MMM-YY	Driver_ID	Quarterly Rating	Rating_Increased
0	2019-01-01	1	2	0
1	2019-01-02	1	2	0
2	2019-01-03	1	2	0
3	2020-01-11	2	1	0

4	2020-01-12	2	1	0
...
19099	2020-01-08	2788	3	0
19100	2020-01-09	2788	3	0
19101	2020-01-10	2788	2	0
19102	2020-01-11	2788	2	0
19103	2020-01-12	2788	2	0

[19104 rows x 4 columns]

```
df.Rating_Increased.value_counts()
```

```
Rating_Increased
```

```
0    17859
```

```
1     1245
```

```
Name: count, dtype: int64
```

```
df['MMM-YY']=pd.to_datetime(df['MMM-YY'],format='%d%m%y')
```

```
df=df.sort_values(by=['Driver_ID','Income'])
```

```
def Income_Increasing(Income):
```

```
    return (Income>Income.shift(1)).astype('int64')
```

```
df['Salary_Increased']=df.groupby('Driver_ID')
```

```
['Income'].apply(Income_Increasing).reset_index(level=0,drop=True)
```

```
print(df[["MMM-YY", "Driver_ID", "Income", "Salary_Increased"]])
```

	MMM-YY	Driver_ID	Income	Salary_Increased
0	2019-01-01	1	57387	0
1	2019-01-02	1	57387	0
2	2019-01-03	1	57387	0
3	2020-01-11	2	67016	0
4	2020-01-12	2	67016	0
...
19099	2020-01-08	2788	70254	0
19100	2020-01-09	2788	70254	0
19101	2020-01-10	2788	70254	0
19102	2020-01-11	2788	70254	0
19103	2020-01-12	2788	70254	0

[19104 rows x 4 columns]

```
df.Salary_Increased.value_counts()
```

```
Salary_Increased
```

```
0    19060
```

```
1         44
```

```
Name: count, dtype: int64
```

Class Imbalance Treatment

```
df.columns
```

```
Index(['Unnamed: 0', 'MMM-YY', 'Driver_ID', 'Age', 'Gender', 'City',  
      'Education_Level', 'Income', 'Dateofjoining',  
      'LastWorkingDate',  
      'Joining Designation', 'Grade', 'Total Business Value',  
      'Quarterly Rating', 'Hasleftcompany', 'Birthyear',  
      'AgeAtJoining',  
      'Tenure', 'Salary_Range', 'Quarterly_Range',  
      'Riders_Age_Category',  
      'Rating_Increased', 'Salary_Increased'],  
      dtype='object')
```

```
df['Hasleftcompany'].isnull().sum()
```

```
0
```

```
df['Hasleftcompany'].value_counts()
```

```
Hasleftcompany
```

```
0    17488
```

```
1     1616
```

```
Name: count, dtype: int64
```

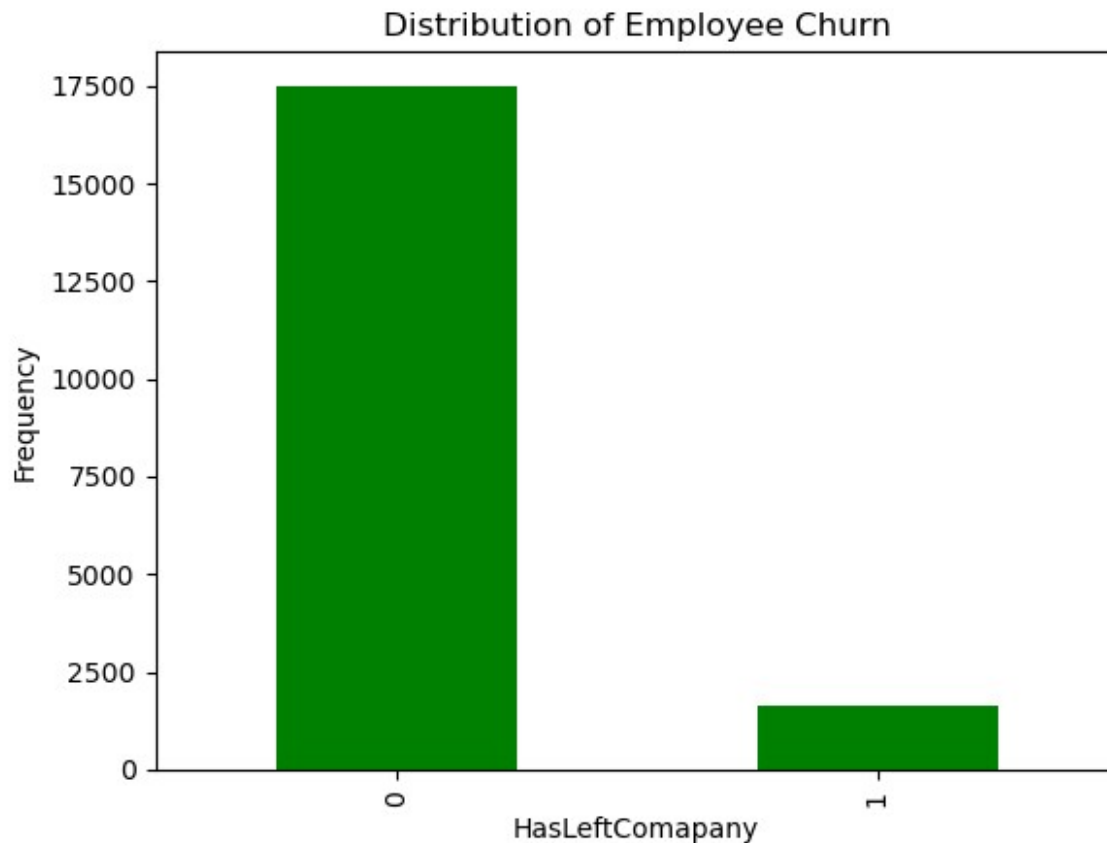
```
df['Hasleftcompany'].value_counts().plot(kind='bar',color='green')
```

```
plt.xlabel("HasLeftComapany")
```

```
plt.ylabel("Frequency")
```


```
plt.title("Distribution of Employee Churn")
```

```
plt.show()
```

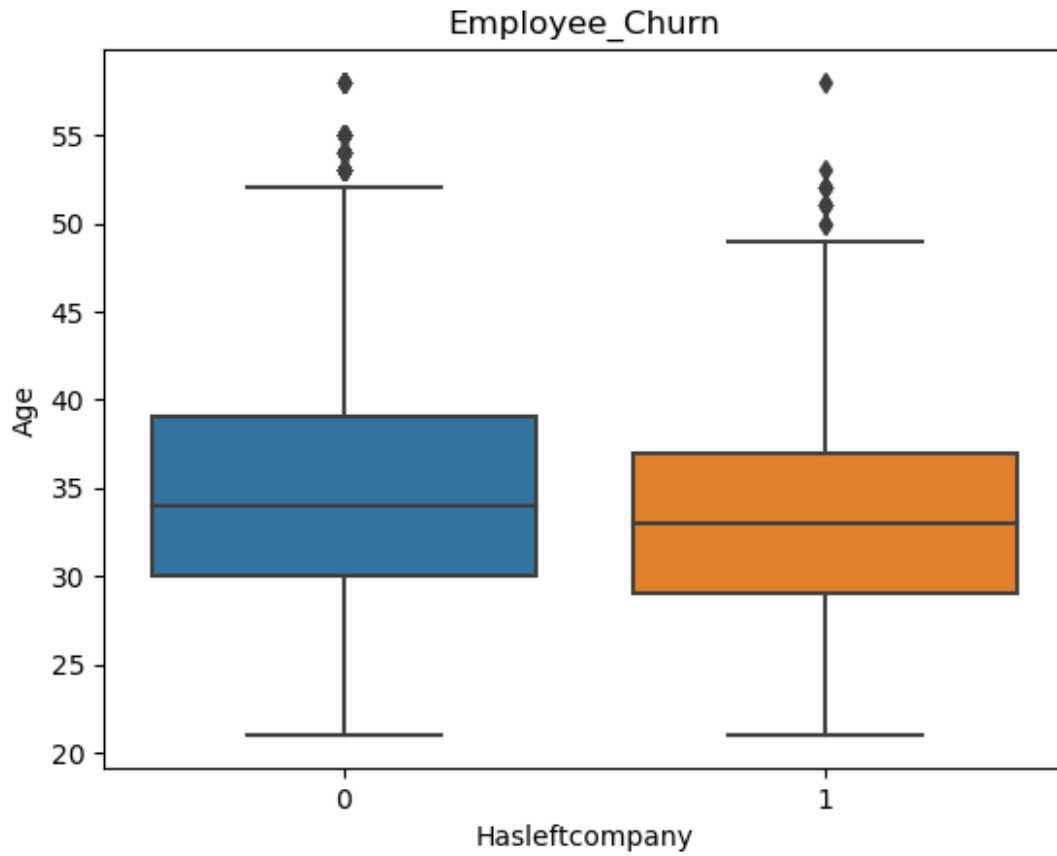


```
numeric_dtypes_df=df.select_dtypes(include=[np.number])

plt.figure(figsize=(12,10))
corr=numeric_dtypes_df.corr()
sns.heatmap(corr,annot=True,cmap='coolwarm',fmt='0.2f',linewidths=0.5,
linecolor='green',cbar_kws={'shrink':0.8},annot_kws={'size':10})
plt.title("Correalation Heatmap", fontsize = 16)
plt.xlabel("Features", fontsize = 14)
plt.ylabel("Features", fontsize = 14)
plt.tight_layout()
plt.show()
```



```
sns.boxplot(x=df['Hasleftcompany'],y=df['Age'],data=df)
plt.title("Employee_Churn")
plt.show()
```




```
Gender_Churn=df.groupby('Gender')['Hasleftcompany'].mean()
print(Gender_Churn)

sns.barplot(x=Gender_Churn.index,y=Gender_Churn.values,color='green',palette='plasma')
plt.title("Churn According to Gender")
plt.xlabel("Gender")
plt.ylabel("Churn")
plt.show()
```

Gender

0.0	0.085202
1.0	0.083740

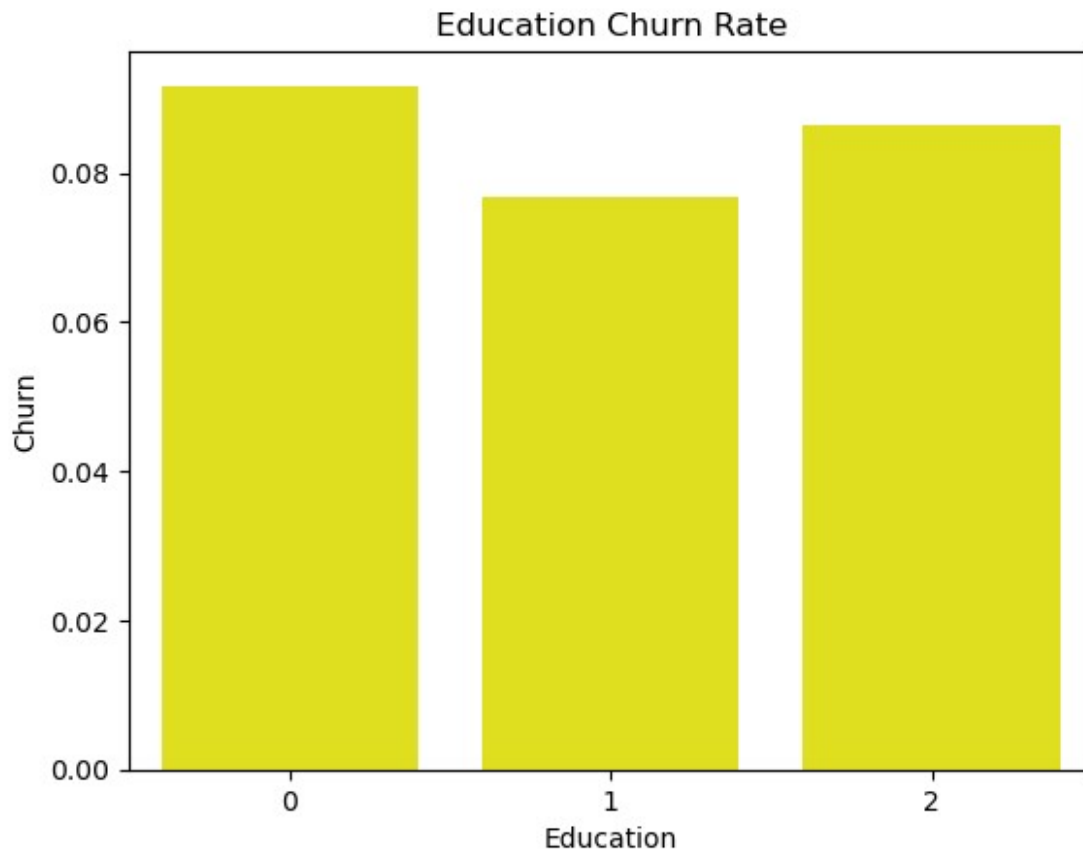
Name: Hasleftcompany, dtype: float64




```
Education_Churn=df.groupby('Education_Level')['Hasleftcompany'].mean()  
print(Education_Churn)
```

```
sns.barplot(x=Education_Churn.index,y=Education_Churn.values,color='yellow')  
plt.title("Education Churn Rate")  
plt.xlabel("Education")  
plt.ylabel("Churn")  
plt.show()
```

```
Education_Level  
0    0.091662  
1    0.076777  
2    0.086455  
Name: Hasleftcompany, dtype: float64
```





```
plt.figure(figsize=(12,6))
sns.histplot(df[df['Hasleftcompany']==0]
['Income'],kde=True,color='blue',label='left')
sns.histplot(df[df['Hasleftcompany']==1]
['Income'],kde=True,color='red',label='Stayed')
plt.xlabel("Income")
plt.ylabel("Frequency")
plt.title("Income Distribution According Churn Status")
plt.legend()
plt.show()
```



```
plt.figure(figsize=(12,6))

sns.violinplot(x='Hasleftcompany',y='Tenure',data=df,color='blue',pale
tte='plasma',bw=0.5,linewidth=1.5,inner='quartile',legend=('HasleftCom
pany'))
plt.xlabel("Churn", fontsize = 14)
plt.ylabel("Tenure", fontsize = 14)
plt.title("Churn Rate By Tenure", fontsize = 16)
plt.xticks([0,1],["Stay", "Left"], fontsize = 12)
plt.grid()
plt.show()
```




```
sns.boxplot(x='Hasleftcompany',y='Quarterly  
Rating',data=df,color='pink')  
plt.title('Quarterly Rating vs Churn')  
plt.show()
```



```

C22    0.061805
C23    0.105948
C24    0.083062
C25    0.092466
C26    0.074799
C27    0.076336
C28    0.086384
C29    0.056667
C3     0.081633
C4     0.089965
C5     0.073171
C6     0.083333
C7     0.085386
C8     0.074438
C9     0.101923
Name: Hasleftcompany, dtype: float64

```



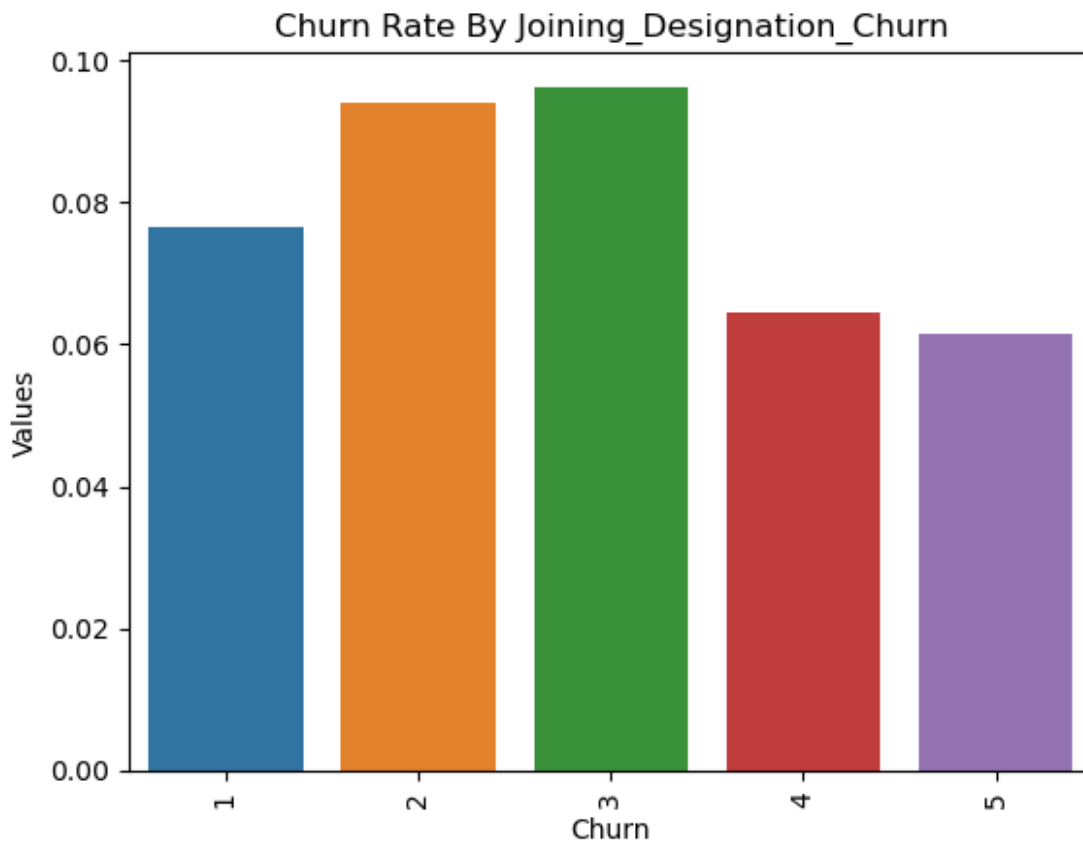
```

Joining_Designation_churn=df.groupby('Joining Designation')
['Hasleftcompany'].mean()
print(Joining_Designation_churn)
sns.barplot(x=Joining_Designation_churn.index,y=Joining_Designation_churn.values)
plt.xlabel("Churn")

```

```
plt.ylabel("Values")
plt.title("Churn Rate By Joining_Designation_Churn")
plt.xticks(rotation = 90)
plt.show()
```

```
Joining Designation
1    0.076493
2    0.094039
3    0.096242
4    0.064516
5    0.061538
Name: Hasleftcompany, dtype: float64
```



Standardization

```
Numerical_Features=df.select_dtypes(include=[np.number])
from sklearn.preprocessing import StandardScaler
Scaler=StandardScaler()
Standarized_data=Scaler.fit_transform(Numerical_Features)
print(Standarized_data)
```

```
[[-1.73196015 -1.74493508 -1.06733399 ... -0.08989762 -0.26403172
-0.04804685]
[-1.73177882 -1.74493508 -1.06733399 ... -0.08989762 -0.26403172
-0.04804685]
[-1.73159749 -1.74493508 -1.06733399 ... -2.45948835 -0.26403172
-0.04804685]
...
[ 1.73159749  1.69290216 -0.74721868 ... -0.6765064  -0.26403172
-0.04804685]
[ 1.73177882  1.69290216 -0.74721868 ... -0.6765064  -0.26403172
-0.04804685]
[ 1.73196015  1.69290216 -0.74721868 ... -0.6765064  -0.26403172
-0.04804685]]
```

Encoding

```
Encoded_df=pd.get_dummies(df[['City','Education_Level','Joining
Designation']],drop_first=True)
```

Encoded_df

	Education_Level	Joining Designation	City_C10	City_C11
City_C12 \				
0	2	1	False	False
False				
1	2	1	False	False
False				
2	2	1	False	False
False				
3	2	2	False	False
False				
4	2	2	False	False
False				
...
...				

19099	2	2	False	False
False				
19100	2	2	False	False
False				
19101	2	2	False	False
False				
19102	2	2	False	False
False				
19103	2	2	False	False
False				

	City_C13	City_C14	City_C15	City_C16	City_C17	...	City_C27
\							
0	False	False	False	False	False	...	False
1	False	False	False	False	False	...	False

2	False	False	False	False	False	...	False
3	False	False	False	False	False	...	False
4	False	False	False	False	False	...	False
...
19099	False	False	False	False	False	...	True
19100	False	False	False	False	False	...	True
19101	False	False	False	False	False	...	True
19102	False	False	False	False	False	...	True
19103	False	False	False	False	False	...	True
	City_C28	City_C29	City_C3	City_C4	City_C5	City_C6	City_C7
\							
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	False	False	False	False	False	True
4	False	False	False	False	False	False	True
...
19099	False	False	False	False	False	False	False
19100	False	False	False	False	False	False	False
19101	False	False	False	False	False	False	False
19102	False	False	False	False	False	False	False
19103	False	False	False	False	False	False	False
	City_C8	City_C9					
0	False	False					
1	False	False					
2	False	False					
3	False	False					
4	False	False					


```
...      ...      ...
19099    False    False
19100    False    False
19101    False    False
19102    False    False
19103    False    False

[19104 rows x 30 columns]
```

Actionable Insights & Recommendations

INSIGHTS

Recommendations

Questions

Data Structure and Overview

```
df.shape
(19104, 23)

df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 19104 entries, 0 to 19103
Data columns (total 23 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   Unnamed: 0                             19104 non-null  int64
 1   MMM-YY                                 19104 non-null  datetime64[ns]
 2   Driver_ID                             19104 non-null  int64
 3   Age                                    19104 non-null  float64
 4   Gender                                 19104 non-null  float64
 5   City                                   19104 non-null  object
 6   Education_Level                       19104 non-null  int64
 7   Income                                 19104 non-null  int64
 8   Dateofjoining                         19104 non-null  datetime64[ns]
 9   LastWorkingDate                       1616 non-null   datetime64[ns]
10   Joining Designation                   19104 non-null  int64
11   Grade                                 19104 non-null  int64
12   Total Business Value                  19104 non-null  int64
13   Quarterly Rating                      19104 non-null  int64
14   Hasleftcompany                       19104 non-null  int64
15   Birthyear                             19104 non-null  int64
16   AgeAtJoining                         19104 non-null  int64
17   Tenure                                19104 non-null  int64
```

```

18 Salary_Range      19104 non-null  category
19 Quarterly_Range   11425 non-null  category
20 Riders_Age_Category 19069 non-null  category
21 Rating_Increased   19104 non-null  int32
22 Salary_Increased   19104 non-null  int64
dtypes: category(3), datetime64[ns](3), float64(2), int32(1),
int64(13), object(1)
memory usage: 3.0+ MB

```

Descriptive Statistics

```

temp = df[["Age", "Income", "Total Business Value", "Quarterly
Rating"]].aggregate([np.mean, np.median, np.std])
temp

```

	Age	Income	Total Business Value	Quarterly
Rating				
mean	34.668435	65652.025126	5.716621e+05	
2.008899				
median	34.000000	60087.000000	2.500000e+05	
2.000000				
std	6.247912	30914.515344	1.128312e+06	
1.009832				

```

df['Driver_ID'].nunique()
2381

```

Temporal Analysis

```

df['Hasleftcompany'].value_counts()
Hasleftcompany
0    17488
1     1616
Name: count, dtype: int64

df['Dateofjoining']=pd.to_datetime(df['Dateofjoining'])
df['Month_Year_Joining']=df['Dateofjoining'].dt.to_period('M')
per_month_jooining_count=df.groupby('Month_Year_Joining')
['Driver_ID'].count()
per_month_jooining_count

```

Month_Year_Joining	
2013-04	31
2013-05	24
2013-06	59
2013-07	63
2013-08	33
...	
2020-08	325

```

2020-09      314
2020-10      139
2020-11       93
2020-12       59
Freq: M, Name: Driver_ID, Length: 85, dtype: int64

Average_Tenure_of_driver=df['Tenure'].mean()
print(f"Average Tenure of Driver is: {Average_Tenure_of_driver:.2f}")

Average Tenure of Driver is: 2307.53

```

Feature Engineering

```

df['Hasleftcompany'].value_counts()

Hasleftcompany
0      17488
1       1616
Name: count, dtype: int64

df['Year_of_Joining']=df['Dateofjoining'].dt.year
df['Year_of_Joining'].value_counts(ascending=False)

Year_of_Joining
2018      4936
2019      4515
2020      3667
2015      1965
2016      1625
2017      1100
2013       693
2014       603
Name: count, dtype: int64

df['Month_of_Joining']=df['Dateofjoining'].dt.month_name()
df['Month_of_Joining'].value_counts(ascending=False)

Month_of_Joining
July      2730
May       2362
October   2095
June      1973
August    1886
November  1867
September 1449
January   1381
December  1261
April     1014
February   684
March      402
Name: count, dtype: int64

```

```
df['Quarter_of_Joining']=df['Dateofjoining'].dt.quarter
df['Quarter_of_Joining'].value_counts(ascending=False)
```

```
Quarter_of_Joining
3      6065
2      5349
4      5223
1      2467
Name: count, dtype: int64
```

```
df['Day_of_Joining']=df['Dateofjoining'].dt.day_name()
df['Day_of_Joining'].value_counts(ascending=False)
```

```
Day_of_Joining
Friday      4439
Sunday      3153
Thursday    2823
Monday      2769
Saturday    2705
Tuesday     2103
Wednesday   1112
Name: count, dtype: int64
```

EDA


```
fig,axes=plt.subplots(1,3,figsize=(15,5))

sns.histplot(x="Age",data=df,kde=True,bins=10,ax=axes[0],color='pink')
axes[0].set_title("Age Distribution")

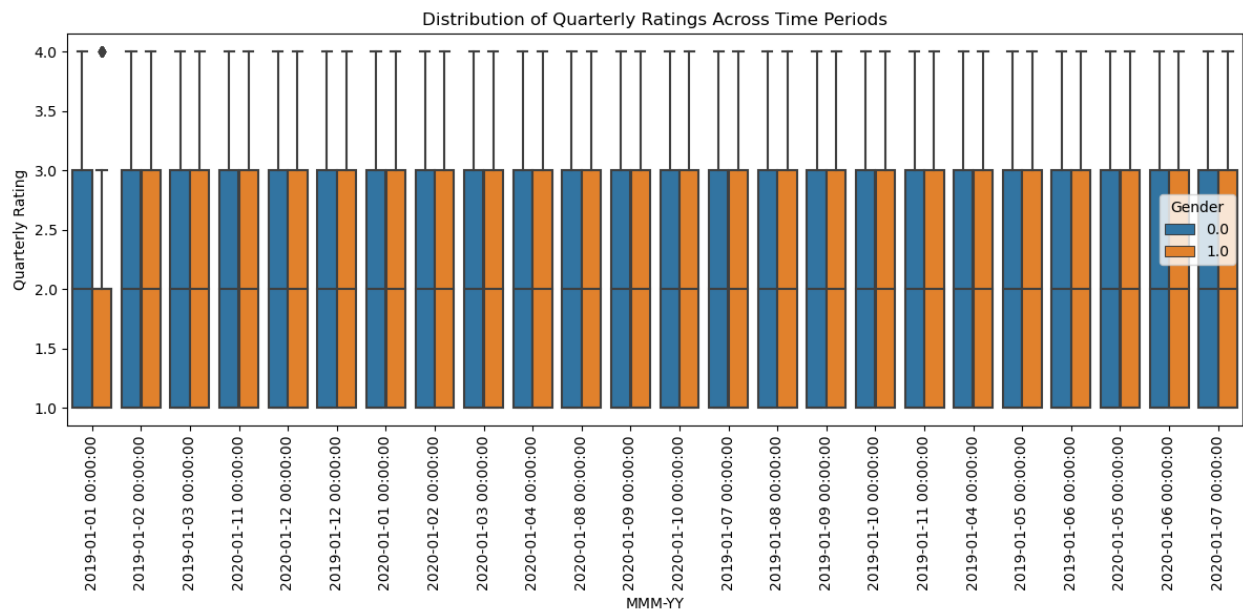
sns.histplot(x='Income',data=df,kde=True,bins=10,ax=axes[1],color='green')
axes[1].set_title("Income Distribution")

sns.histplot(x='Total Business Value',data=df,kde=True,bins=10,ax=axes[2],color='black')
axes[2].set_title("Total Buisness Distribution")

plt.tight_layout()
plt.show()
```




```
plt.figure(figsize=(12,6))
sns.boxplot(x='MMM-YY',y='Quarterly Rating',hue='Gender',data=df)
plt.title('Distribution of Quarterly Ratings Across Time Periods')
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()
```



```
plt.figure(figsize=(14, 6))
sns.boxplot(x='Driver_ID', y='Quarterly Rating', hue='Gender',
data=df_filtered, palette="Set2")
plt.xticks(rotation=90)
plt.title('Distribution of Quarterly Ratings by Drivers (Grouped by Gender)')
plt.xlabel('Driver ID')
plt.ylabel('Quarterly Rating')
plt.legend(title='Gender', loc='upper right')
```

```
plt.tight_layout()
plt.show()
```

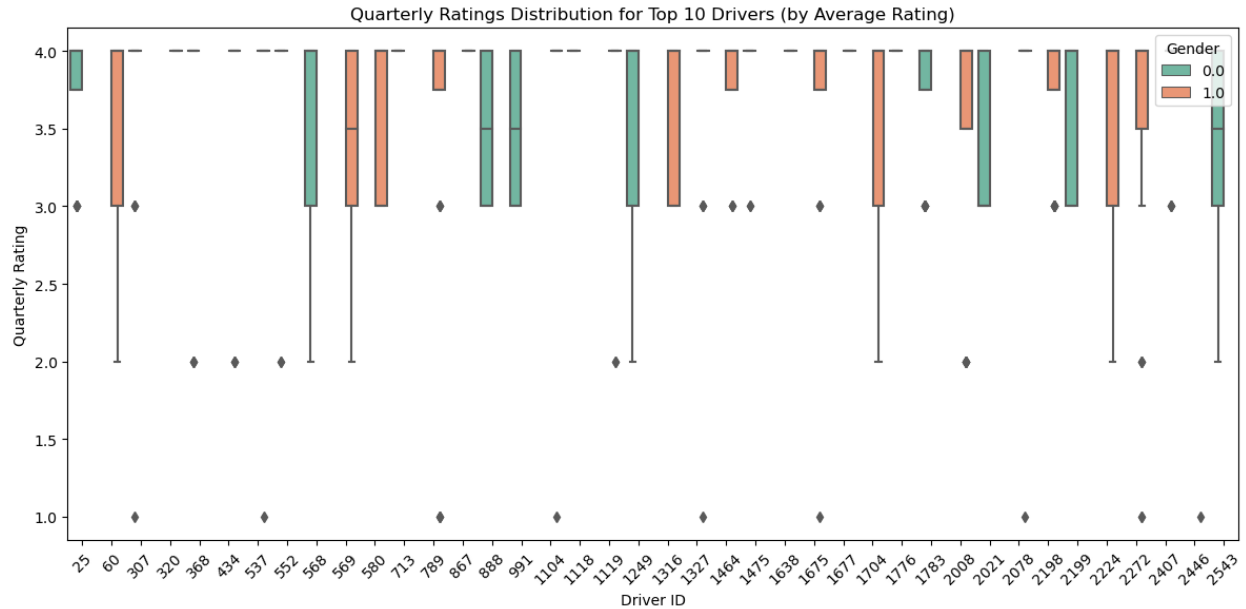


```
top_drivers = (
    df.groupby('Driver_ID')['Quarterly Rating']
      .mean()
      .sort_values(ascending=False)
      .head(40) # Select top 10 drivers
      .index
)
df_top = df[df['Driver_ID'].isin(top_drivers)]

# Step 3: Create a boxplot for the top 10 drivers
plt.figure(figsize=(12, 6))
sns.boxplot(x='Driver_ID', y='Quarterly Rating', hue='Gender',
            data=df_top, palette="Set2")

# Customize the plot
plt.title('Quarterly Ratings Distribution for Top 10 Drivers (by Average Rating)')
plt.xlabel('Driver ID')
plt.ylabel('Quarterly Rating')
plt.xticks(rotation=45)
plt.legend(title='Gender', loc='upper right')
plt.tight_layout()

# Display the plot
plt.show()
```


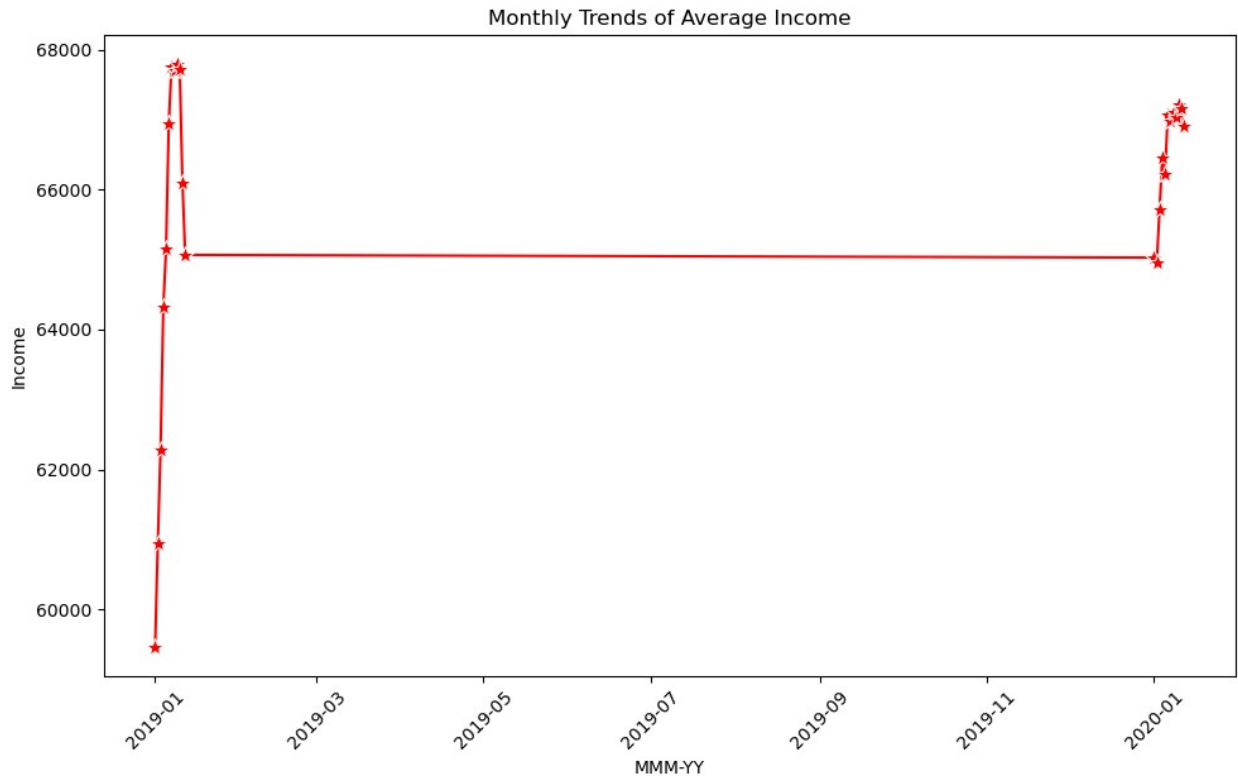


```
plt.figure(figsize=(10,6))
monthly_income=df.groupby('MMM-YY')['Income'].mean().reset_index()

sns.lineplot(x='MMM-YY',y='Income',data=monthly_income,marker='*',color='red',markersize=10)
plt.title("Monthly Trends of Average Income")
plt.tight_layout()
plt.xticks(rotation = 45)
plt.show()

plt.figure(figsize=(10,6))
monthly_income=df.groupby('MMM-YY')['Total Business Value'].mean().reset_index()

sns.lineplot(x='MMM-YY',y='Total Business Value',data=monthly_income,marker='*',color='red',markersize=10)
plt.title("Monthly Trends of Average Income")
plt.tight_layout()
plt.xticks(rotation = 45)
plt.show()
```



Missing Values Handling

```
no_null_Lastworking_date=df.LastWorkingDate.dropna().values
df['LastWorkingDate']=df['LastWorkingDate'].fillna(pd.to_datetime('today').normalize())
```

```
df.LastWorkingDate.isnull().sum()
```

0

Correlation and Relationships

```
Correlation_age_income=df['Age'].corr(df['Income'])
Correlation_age_income
```

0.19084585445978905

```
Education_status=df.groupby('Education_Level')['Total Business Value'].describe()
print("Statics How the Educaton_Level affecting the Total Business Value")
Education_status
```

Statics How the Educaton_Level affecting the Total Business Value

	count	mean	std	min
25% \ Education_Level				
0	5913.0	565410.657872	1.092937e+06	-2628700.0 0.0
1	6864.0	601287.867133	1.227469e+06	-5483890.0 0.0
2	6327.0	545364.175755	1.044904e+06	-6000000.0 0.0

	50%	75%	max
Education_Level			
0	239180.0	689660.0	23550000.0
1	270885.0	721917.5	33747720.0
2	246450.0	676765.0	17651940.0

```
Education_status=df.groupby('City')['Total Business Value'].describe()
print("Statics How the City affecting the Total Business Value")
Education_status
```

Statics How the City affecting the Total Business Value

	count	mean	std	min	25%	50%
75% \ City						

C1	677.0	531560.280650	1.028461e+06	-442150.0	0.0	263930.0
602580.0						
C10	744.0	540753.736559	9.124536e+05	-500000.0	0.0	251815.0
676185.0						
C11	468.0	538549.145299	1.099731e+06	-439300.0	0.0	203935.0
627980.0						
C12	727.0	667282.310867	1.246261e+06	-484900.0	0.0	299930.0
848230.0						
C13	569.0	796263.075571	2.160790e+06	-1304840.0	0.0	266330.0
794250.0						
C14	648.0	607931.635802	1.321692e+06	-582010.0	0.0	203950.0
743022.5						
C15	761.0	553266.636005	9.613450e+05	-250000.0	0.0	250000.0
705790.0						
C16	709.0	632585.712271	1.186105e+06	-1477940.0	0.0	271470.0
798670.0						
C17	440.0	429160.204545	9.620125e+05	-1496650.0	0.0	157025.0
504522.5						
C18	544.0	550106.250000	1.112772e+06	-6000000.0	0.0	300280.0
713742.5						
C19	579.0	630978.151986	1.393624e+06	-1704230.0	0.0	301750.0
749550.0						
C2	472.0	553365.084746	1.147130e+06	-232800.0	0.0	201305.0
565440.0						
C20	1008.0	468535.605159	9.012326e+05	-3791250.0	0.0	153470.0
614690.0						
C21	603.0	572684.776119	9.606052e+05	-1629620.0	0.0	250000.0
705390.0						
C22	809.0	559749.431397	9.485263e+05	-1850000.0	0.0	300000.0
707770.0						
C23	538.0	423986.561338	7.554442e+05	-665480.0	0.0	151080.0
499980.0						
C24	614.0	584712.426710	1.115906e+06	-647520.0	0.0	252000.0
654282.5						
C25	584.0	507575.119863	9.842604e+05	-500000.0	0.0	206200.0
571935.0						
C26	869.0	661837.445339	1.386692e+06	-5483890.0	0.0	276610.0
796830.0						
C27	786.0	572039.312977	9.689232e+05	-2628700.0	0.0	321580.0
752167.5						
C28	683.0	591406.778917	1.144723e+06	-2910060.0	0.0	250000.0
765605.0						
C29	900.0	736637.511111	1.332774e+06	-619680.0	0.0	369300.0
854320.0						
C3	637.0	458003.940345	7.576680e+05	-1590270.0	0.0	242160.0
574090.0						
C4	578.0	556092.266436	1.004997e+06	-500000.0	0.0	203515.0
669765.0						
C5	656.0	634855.975610	1.271910e+06	-831520.0	0.0	250000.0

```

750000.0
C6      660.0   566042.954545   1.033945e+06   -921510.0   0.0   300000.0
686437.5
C7      609.0   484569.228243   8.848838e+05   -344010.0   0.0   200000.0
600000.0
C8      712.0   566328.539326   9.622600e+05   -154360.0   0.0   273900.0
756330.0
C9      520.0   467914.865385   8.592157e+05   -242830.0   0.0   241755.0
558292.5

```

```

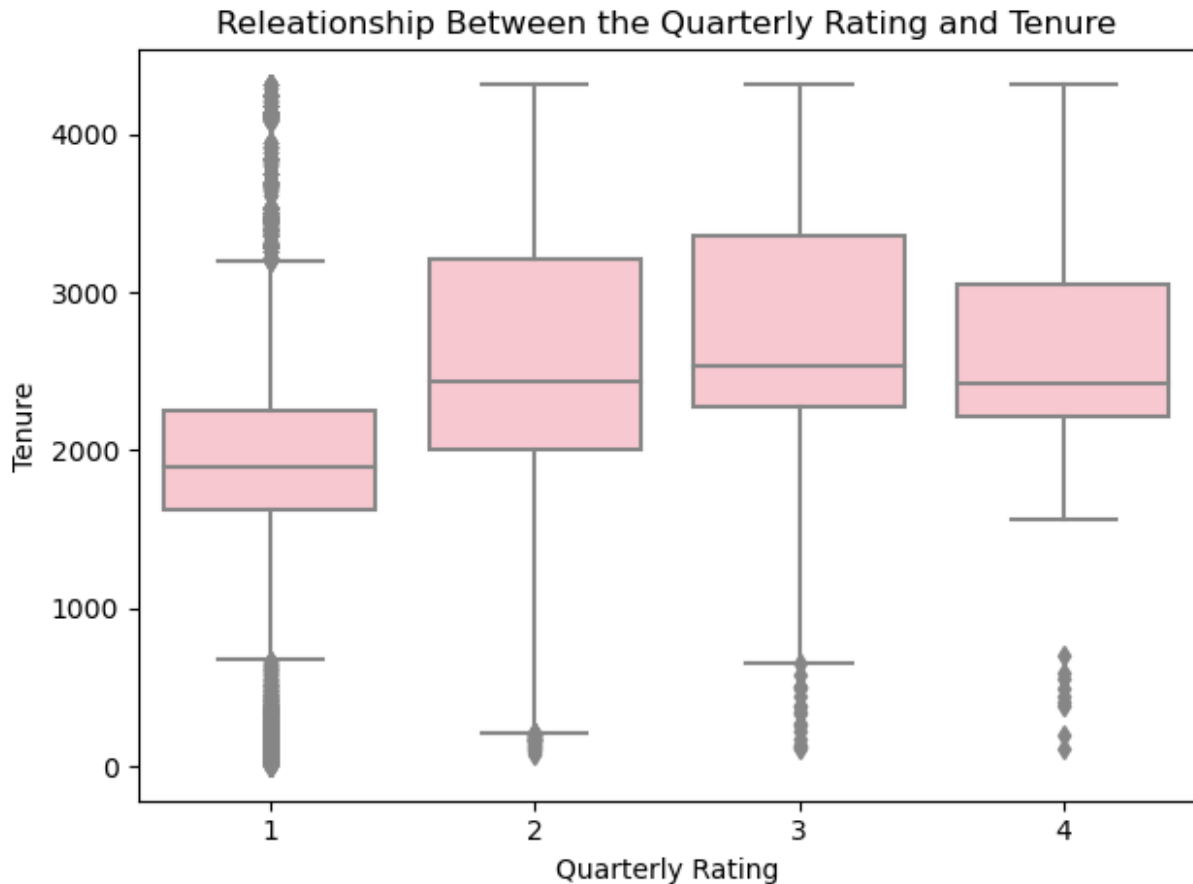
max
City
C1      13737020.0
C10     10703580.0
C11     13435570.0
C12     16873160.0
C13     33747720.0
C14     17651940.0
C15     12506660.0
C16     11454170.0
C17     12836130.0
C18     12849540.0
C19     23550000.0
C2      10951600.0
C20      7799990.0
C21      9000000.0
C22     13588750.0
C23      6350000.0
C24     11040770.0
C25     15059230.0
C26     16979740.0
C27     10305800.0
C28     12200000.0
C29     16606860.0
C3       7435230.0
C4       8560000.0
C5     15993610.0
C6     11894240.0
C7       9100910.0
C8     12941160.0
C9     12701500.0

```

```

sns.boxplot(x='Quarterly Rating',y='Tenure',data=df,color='pink')
plt.title("Releationship Between the Quarterly Rating and Tenure")
plt.tight_layout()
plt.show()

```



```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score,
confusion_matrix

df=df.drop(columns=['Unnamed: 0'])

label_encoder=LabelEncoder()
df['Gender']=label_encoder.fit_transform(df['Gender'])
df['City']=label_encoder.fit_transform(df['City'])
df['Education_Level'] =
label_encoder.fit_transform(df['Education_Level'])
df['Joining Designation'] = label_encoder.fit_transform(df['Joining
Designation'])
df['Grade'] = label_encoder.fit_transform(df['Grade'])
df['Salary_Range'] = label_encoder.fit_transform(df['Salary_Range'])
df['Quarterly_Range'] =
label_encoder.fit_transform(df['Quarterly_Range'])
df['Riders_Age_Category'] =
label_encoder.fit_transform(df['Riders_Age_Category'])

```

```

X=df.drop(columns=['Hasleftcompany', 'Dateofjoining',
'LastWorkingDate', 'MMM-
YY', 'Month_Year_Joining', 'Month_of_Joining', 'Day_of_Joining'])
y=df['Hasleftcompany']

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)

scaler=StandardScaler()
X_train_scaled=scaler.fit_transform(X_train)
X_test_scaled=scaler.transform(X_test)

model=RandomForestClassifier(random_state=42)
model.fit(X_train_scaled,y_train)

y_pred=model.predict(X_test_scaled)

print('Accuracy',accuracy_score(y_test,y_pred))
print('Classification Report:')
print(classification_report(y_test,y_pred))
print('Confusion Matrix:')
conf_matrix=confusion_matrix(y_test,y_pred)

sns.heatmap(conf_matrix,annot=True,fmt='d',cmap='Blues',xticklabels=['Stayed','Left'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

features_Importance=model.feature_importances_
feature_name=X.columns
sorted_idx=features_Importance.argsort()

plt.barh(feature_name[sorted_idx],features_Importance[sorted_idx])
plt.xlabel("Feature Importance")
plt.title("Feature Importance - RandomForestClassifier")
plt.show()


```


Accuracy 0.9965977492802931

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3490
1	1.00	0.96	0.98	331
accuracy			1.00	3821
macro avg	1.00	0.98	0.99	3821
weighted avg	1.00	1.00	1.00	3821

Confusion Matrix:





Recommendations

Actionable Insights & Recommendations

```
import os
```

```
df.to_csv('Cleaned_Dataset.csv', index=False)
```

https://app.powerbi.com/links/IfRjsPKjgS?ctid=ed77d40f-8c11-413d-96e2-467edfd73d60&pbi_source=linkShare